

# Computer Architecture

Ali Muhammad aa07190 — Section L3

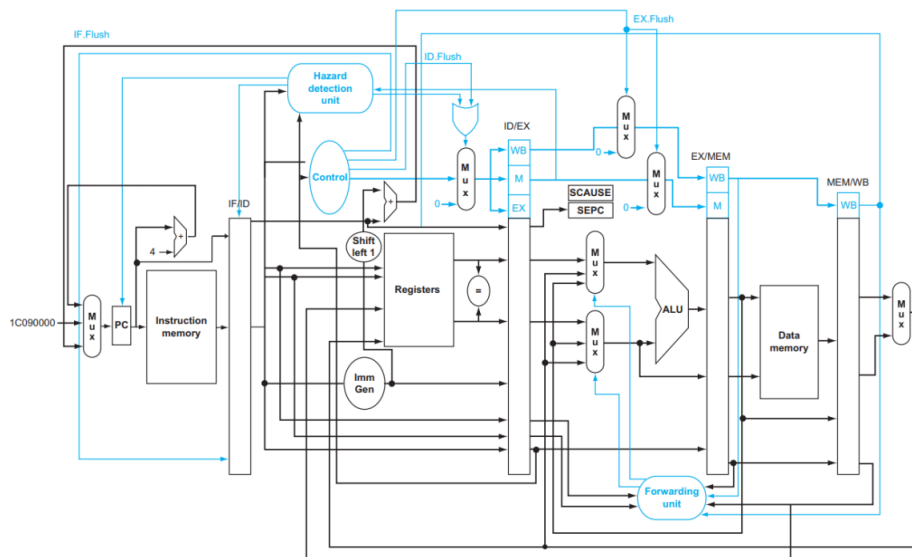
Iqra Ahmed ia07674 — Section L7

## Homework 3

Question:	1	2	3	4	5	6	7	8	9	10	11	Total
Points:	5	5	5	5	5	5	50	5	5	5	5	100
Score:												

### Question 1: Pipeline Register Design [5 marks]

What are the total number of bits needed for each pipeline register in the following pipelined RISC-V processor. Also break it down into the individual fields of each pipeline register and how many bits are needed for each field.



**Solution:**

The RISC-V Processor can be broken down into the following pipelines:

1. **IF / ID - Instruction Fetch / Instruction Decode:** Program Counter (64 bits) and Instruction Memory that gives 32 bit instruction.

Total bits =  $64 + 32 = 96$  bits.

2. **ID / EX - Instruction Decode / Execute**

Program Counter - 64 bits.

From Register File: ReadData1 (64 bits), ReadData2 (64 bits), RS1 (5 bits), RS2 (5 bits), RD (5 bits).

Immediate Generator - 64 bits.

ALUOp 4 bits, ALUSrc 1 bit

Branch, MemRead, MemWrite, MemtoReg, RegWrite each of 1 bit (5 bits in total)

Total bits =  $64 + 64 + 64 + 5 + 5 + 5 + 64 + 4 + 1 + 5 = 281$  bits.

3. **EX / MEM - Execute / Access operand from Data Memory:**

ALU Result - 64 bit.

Mux Out (ReadData2 or Immediate Value) - 64 bit.

RD - 5 bits.

Left Shifted Immediate Value - 64 bit

Branch, MemRead, MemWrite, RegWrite, MemtoReg each of 1 bit (5 bits in total)

Total bits =  $64 + 64 + 64 + 5 + 5 = 202$  bits.

4. **Mem / WB - Memory / Write Back:**

Read Data from Data Memory - 64 bits.

ALU Result - 64 bits.

RD - 5 bits

MemtoReg and RegWrite each of 1 bit (2 bits in total)

Total bits =  $64 + 64 + 5 + 2 = 135$  bits.

Total number of bits for each pipeline register in RISC-V pipelined processor =  $96 + 281 + 202 + 135 = 714$  bits.

[The image given above is very vague without proper labeling, hence influence was taken from the pipelined RISC-V processor diagram given with the project details file.]

**Question 2: Pipeline Basics [5 marks]**

In this exercise we examine how pipelining affects the clock cycle time of the processor. Problems in this exercise assume that individual stages of the datapath have the following latencies:

IF	ID	EX	MEM	WB
200ps	300ps	150ps	350ps	250ps

Also, assume that instructions executed by the processor are broken down as follows:

R-Type	beq	ld	sd
40%	15%	20%	25%

- (i) What is the clock cycle time in a pipelined and non-pipelined processor?

**Solution:** In pipelining, all stages must adhere to a single clock cycle that should be enough to accomodate the operation that takes the longest time. Therefore, for pipelined processor, clock cycle time should be 350ps.

In non pipelining, each instruction would go through all stages and then new instruction would be executed. This happens in a single clock cycle, therefore, for a non-pipelined processor clock cycle time =  $200 + 300 + 150 + 350 + 250 = 1250\text{ps}$ .

- (ii) What is the total latency of an ld instruction in a pipelined and non-pipelined processor?

**Solution:** An ld instruction takes all individual stages of a processor.

Pipelined: Each stage would take 350ps since that is the clock cycle time. Total Latency =  $350 * 5 = 1750\text{ps}$ .

Non-Pipelined: Each stage would take its own time. Total Latency =  $200 + 300 + 150 + 350 + 250 = 1250\text{ps}$ .

- (iii) If we can split one stage of the pipelined datapath into two new stages, each with half the latency of the original stage, which stage would you split and what is the new clock cycle time of the processor?

**Solution:** Since the clock cycle time is based on the stage with the highest latency, it would be ideal to split the stage that takes the longest time. In our case that would be the MEM stage.

The new clock cycle time would then be the stage with the highest latency; ID - Instruction Decode with 300ps.

- (iv) Assuming there are no stalls or hazards, what is the utilization of the data memory?

**Solution:** Data Memory is only used by either ld, or sd instruction that take upto 20% and 25% of the processor. So utilization of the data memory is  $20\% + 25\% = 45\%$  of the clock cycles.

- (v) Assuming there are no stalls or hazards, what is the utilization of the write-register port of the “Registers” unit?

**Solution:** The write register is used by either R-Type instructions that require arithmetic operations or ld instruction as they require the data to be written. So utilization of write register = 40% + 20% = 60% of the clock cycle time.

### Question 3: Pipeline Hazard Basics [5 marks]

A hypothetical processor has 9 stages of a pipeline as shown in table below. The first row in the table below shows the pipeline stage number, second row gives the name of each stage, and third row gives the delay of each stage in Nano-seconds. The name of each stage describes the task performed by it. Each stage takes 1 cycle to execute. This processor stores all the register contents in a compressed fashion. After fetching the operands the operands are first decompressed, and before saving the results in register file, the results are first compressed.

1	2	3	4	5	6	7	8	9
Instruction Fetch	Instruction Decode	Operand Fetch	Decompress Operands	Instruction Execute 1	Instruction Execute 2	Memory Access	Compress results	Register Write back
1ns	1.4ns	1.1ns	2.2ns	2.3ns	2.3ns	1.3ns	2.2ns	1.2ns

- (i) How many cycles are required to implement/execute one instruction on this pipeline?

**Solution:** Since this is a 9 stage pipelined design, it'll take 9 clock cycles.

- (ii) How many cycles are required to execute 17 instructions on this pipeline? Assume that no stall cycles occur during the execution of all instructions.

**Solution:** 17 instructions would take 17 cycles, plus the remaining stages left for the last cycle [overlaps hence can be considered like this.]  
 Total cycles =  $17 + (9 - 1) = 25$  cycles.

- (iii) Assume that all necessary bypass circuitry is implemented in this 9 stage pipeline. How many cycles will the pipeline stall during the execution of below given two instructions?
- $x5 = \text{Load from memory}$
  - $x2 = x5 + x7$

**Solution:** There will be 2 stalls assuming that all necessary bypass circuitry is implemented as when the instruction is loaded from memory in the Memory Access Stage, then it can be bypassed/forwarded to the Instruction Execute 1 stage which results in only **2 stalls**.

- (iv) Assume that no bypass circuitry is implemented in this 9 stage pipeline. How many cycles will the pipeline stall during the execution of above mentioned two instructions?

**Solution:** It will have an additional 2 stalls, as there is no bypass, then we have to wait for register write back, after which operand is fetched and decompressed before execution. So **4 stalls**.

- (v) Assume that all necessary bypass circuitry implemented in this 9 stage pipeline. How many cycles will the pipeline stall during the execution of below given two instructions?
- $x5 = x1 + x3$
  - $x2 = x5 + x7$

**Solution:** There should be 2 stalls again by similar logic as part (iii).

- (vi) Assume that no bypass circuitry is implemented in this 9 stage pipeline. How many cycles will the pipeline stall during the execution of above mentioned two instructions?

**Solution:** 4 stalls by similar logic as part (iv).

- (vii) How much total time (in ns) is required to execute one entire instruction if the nine stages were **not pipelined**?

**Solution:** Total time =  $1 + 1.4 + 1.1 + 2.2 + 2.3 + 2.3 + 1.3 + 2.2 + 1.2 = 15.0\text{ns}$

- (viii) What is the delay of 1 cycle (in ns) when all the nine stages are pipelined?

**Solution:** Delay should be of the instruction with highest latency;  $2.3\text{ns}$

- (ix) What is the total time (in ns) required to execute one entire instruction if the nine stages are pipelined?

**Solution:** Total time = time for 1 cycle  $\times$  stages =  $2.3 \times 9 = 20.7\text{ns}$

(x) Below are given two set of codes. For each set of code, mention if forwarding circuit can avoid all the stalls in the code?

- a.    add   x1, x2, x3  
       add   x6, x7, x8  
       add   x4, x1, x5
- b.    ld     x1, 0(x3)  
       add   x2, x1, x3

**Solution:**

- a. Stalls can be avoided by forwarding
- b. Stalls cannot be avoided by forwarding

**Question 4: Pipeline Diagram [5 marks]**

Consider the following loop.

```

LOOP:  ld x11, 8(x13)
       ld x10, 0(x13)
       add x12, x10, x11
       addi x13, x13, -16
       bne x12, x0, LOOP
  
```

Assume that perfect branch prediction is used (no stalls due to control hazards), that the pipeline has full forwarding support, and that branches are resolved in the EX (as opposed to the ID) stage.

- (i) Show a pipeline execution diagram for the first two iterations of this loop.
- (ii) Mark pipeline stages that do not perform useful work. How often while the pipeline is full do we have a cycle in which all five pipeline stages are doing useful work?

**Solution:** (i) and (ii) [Zoom in to see a better picture, can be zoomed in.]

Instruction \ Cycle	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
ld x11 8(x13)	IF	ID	EX	MEM	WB											
ld x10 0(x13)		IF	ID	EX	MEM	WB										
add x12, x10, x11			IF	ID	stall	EX	MEM	WB								
addi x13, x13, -16				IF	stall	ID	EX	MEM	WB							
bne x12, x0, LOOP				stall	IF	ID	EX	MEM	WB							
ld x11 8(x13)						IF	ID	EX	MEM	WB						
ld x10 0(x13)							IF	ID	EX	MEM	WB					
add x12, x10, x11								IF	ID	stall	EX	MEM	WB			
addi x13, x13, -16									IF	stall	ID	EX	MEM	WB		
bne x12, x0, LOOP										stall	IF	ID	EX	MEM	WB	

The dark blue arrows show where forwarding has been done. The yellow boxes show where there is a stall and the red boxes show stages in the pipeline where no useful work is being performed. As it can be seen clearly that we do not have any cycle in the pipeline during which all five stages are performing useful work.

**Question 5: Energy Consideration in Pipeline Design [5 marks]**

This exercise explores energy efficiency and its relationship with performance. Problems in this exercise assume the following energy consumption for activity in Instruction memory, Registers, and Data memory. You can assume that the other components of the datapath consume a negligible amount of energy. (“Register Read” and “Register Write” refer to the register file only.)

I-Mem	1 Register Read	Register Write	D-Mem Read	D-Mem Write
120pJ	80pJ	60pJ	120pJ	100pJ

Keep in mind that reading two registers (instead of one) will double the energy consumption of register reading. Assume that components in the datapath have the following latencies. You can assume that the other components of the datapath have negligible latencies.

I-Mem	Control	RegisterReadorWrite	ALU	D-MemReadorWrite
200ps	150ps	90ps	90ps	250ps

- (i) How much energy is spent to execute an add instruction in a single-cycle design and in the five-stage pipelined design?

**Solution:** The energy spent will remain the same for a single add instruction in both the single-cycle and five-staged pipelined design. I-Mem, 2 register reads, and 1 register write. Energy =  $120 + 2(80) + 60 = 340\text{pJ}$ .

- (ii) What is the worst-case RISC-V instruction in terms of energy consumption? What is the energy spent to execute it?

**Solution:** The worst instruction would be that takes the most energy in total. We either have two read registers at most, such as the add instruction, however, it does not utilise all the components. Now consider the load (ld) instruction that has a register read and a register write. Energy for ld =  $120 + 2(80) + 60 + 120 = 460\text{pJ}$ . So the worst instruction would be the load instruction taking 460pJ of energy for execution.

- (iii) If energy reduction is paramount, how would you change the pipelined design? What is the percentage reduction in the energy spent by an ld instruction after this change?

**Solution:** I-Mem has to be done for all instructions, and so does Register Read. I-Mem can't be modified, however, we can modify the Register Read such that it takes 2 signals; RegRead1 and RegRead2. Then whatever signal will be activated by the control unit (depending on the instruction), only that register will be read. This way, both registers will not have to be read for every instruction, such as the load instruction. We do not modify Register Write, D-Mem Read or D-Mem Write, as they are only performed when signal is active for those operations for those instructions.

Now ld instruction takes energy =  $120 + 80 + 60 + 120 = 380\text{pJ}$ .

Percentage reduction in energy =  $\frac{80}{460} \times 100 = 17.4\%$

- (iv) What other instructions can potentially benefit from the change discussed in (iii)?

**Solution:** All I-Type instructions would benefit from this such as `addi` as they only require one register read. `jal` instruction would also benefit as it does not read any registers.

- (v) How do your changes from (iii) affect the performance of a pipelined CPU?

**Solution:** Previously, the Control Unit would decode the instructions while registers were being read simultaneously. Now the Control Unit would have to decode the instruction before in order to decide which registers to read. Hence, the Control Unit and Register Reads cannot be done at the same time, and would require separate stages or clock cycles. This would increase the latency of the ID stage (Instruction Decode) as to accommodate the Control Unit as well and could potentially increase the clock cycle time of the processor [although in this question this is not the case, but it can be in some cases].

- (vi) We can eliminate the MemRead control signal and have the data memory be read in every cycle, i.e., we can permanently have MemRead=1. Explain why the processor still functions correctly after this change. If 25% of instructions are loads, what is the effect of this change on clock frequency and energy consumption?

**Solution:** If MemRead control signal is always set to 1, then either it is needed by some instruction such as the load instruction, or it is not needed such as the add instruction in which case it is ignored in the MUX after the data memory as it is not needed and MUX does not select it, so is not written anywhere. So the processor still functions correctly.

If 25% instructions are load, then remaining 75% instructions are not load. There is no effect on the clock frequency, however, the energy consumption would most certainly increase due to the number of unnecessary Memory Reads on the remaining 75% instructions. So increase in energy consumption would be of 120pJ in the remaining 75% instructions of clock cycle time of 250ps. Energy increase =  $\frac{120 \times 75\%}{250} = 0.36$



**Question 6: Pipeline Design Optimization for Memory Access [5 marks]**

ld is the instruction with the longest latency on the RISC-V implementation of single-cycle non-pipelined processor discussed in class. If we modified ld and sd so that there was no offset (i.e., the address to be loaded from/stored to must be calculated and placed in rs1 before calling ld/sd), then no instruction would use both the ALU and Data memory. This would allow us to reduce the clock cycle time. However, it would also increase the number of instructions, because many ld and sd instructions would need to be replaced with ld/add or sd/add combinations.

- (i) What would the new clock cycle time be for non-pipelined processor?

**Solution:** It would be the sum of the latencies of all individual stages of the processor; 600ps [from Dr. Soorat's slides.]

Since the question is vague, I took latencies from slides of Dr. Soorat which had 800ps total latency, and 200ps latency for the Execution stage. So without the Execution stage, the clock cycle time for non-pipelined processor would be 600ps.

- (ii) In the non-pipelined version, would a program with the instruction mix presented in Question 2 run faster or slower on this new CPU? By how much? (For simplicity, assume every ld and sd instruction is replaced with a sequence of two instruction.)

**Solution:**  $\text{CPU Time}_{\text{Before}} = 200 + 300 + 150 + 350 + 250 = 1250ps$

Load and Store instructions comprise of 45% of instructions.

$\text{CPU Time}_{\text{New}} = 1.45 * (200 + 300 + 350 + 250) = 1.45 * 1100 = 1595ps.$

$\text{Speedup} = \frac{1250}{1595} = 0.784.$

The new CPU will run 0.784 times slower.

- (iii) What is the primary factor that influences whether a program will run faster or slower on the new CPU?

**Solution:** The primary factor that influences the speed of a program are the number of load and store instructions. In addition, the way load and store instructions are made can also have an effect; if load and store have only few different addresses, then it may also run somewhat faster.

- (iv) Do you consider the original CPU a better overall design; or do you consider the new CPU a better overall design? Why?

**Solution:** I consider the original CPU a better overall design. Even though the old CPU design has a longer clock cycle, it still has much fewer instructions. Generally, load and store instructions are the primary factors that influence the speed of the program. If load and store are further broken down along with add instructions to calculate offset, this would increase the number of instructions and complexity of the design as now the offset would have to be calculated by the add instruction, and then

stored later on for the load and store instructions to use, so more hardware components or signals would have to be used accordingly. This would undoubtedly increase the complexity of the processor. Moreover, the increase in number of instructions might also cause the processor to perform even slower.

- (v) As a result of the change, the MEM and EX stages of the pipelined version of the processor can be overlapped and the pipeline has only four stages. How will the reduction in pipeline depth affect the cycle time?

**Solution:** The reduction will only effect cycle time if the stage with the longest stage has been affected such that its cycle time is reduced. Since they are only overlapped and not changed in any way, the clock cycle time will not be changed as longest latency time would still remain same.

- (vi) How might this change improve the performance of the pipeline?

**Solution:** For lesser instructions with fewer load and store instructions, the new CPU would perform better.

- (vii) How might this change degrade the performance of the pipeline?

**Solution:** For more number of instructions with higher number of load and store instructions, the new CPU would perform poorly as compared to the old CPU design.

**Question 7: Pipeline Design for New Instructions [50 marks]**

Consider the following instructions that are not found in the RISC-V architecture:

- i. **Load Word Register**  
lwr rd, rs2(rs1) //  $\text{Reg}[\text{rd}] = \text{Mem}[\text{Reg}[\text{rs1}] + \text{Reg}[\text{rs2}]]$
- ii. **Add 3 Operands**  
add3 rd, rs1, rs2, rs3 //  $\text{Reg}[\text{rd}] = \text{Reg}[\text{rs1}] + \text{Reg}[\text{rs2}] + \text{Reg}[\text{rs3}]$
- iii. **Add to Memory**  
addm rd, Offset(rs) //  $\text{Reg}[\text{rd}] = \text{Reg}[\text{rd}] + \text{Mem}[\text{Offset} + \text{Reg}[\text{rs}]]$
- iv. **Branch Equal to Memory**  
beqm rs1, offset(rs2), rs3 // if( $\text{Reg}[\text{rs1}] == \text{Mem}[\text{Offset} + \text{Reg}[\text{rs2}]]$ )  
//  $\text{PC} = \text{PC} + \text{Reg}[\text{rs3}]$
- v. **Store Word and Increment**  
swinc rs2, offset(rs1) //  $\text{Mem}[\text{Reg}[\text{rs1}] + \text{offset}] = \text{Reg}[\text{rs2}]$ ,  
//  $\text{Reg}[\text{rs1}] = \text{Reg}[\text{rs1}] + 4$

We want to modify the RISC-V processor to support the above instructions. For parts (b) and (d) below, you can use a printed version of the figures in the book over which you can draw your suggested modifications (no need to draw the entire diagram from scratch).

For each of the above instructions, do the following:

- (a) (1 mark) Suggest if any of the existing instruction formats is a good choice to encode the new instruction. If not, then propose a new instruction format.
- (b) (3 marks) Modify the datapath and control signals of the single-cycle RISC-V processor (Figure 4.17 of the book) to execute the new instruction using the instruction format suggested in part (a). Use the minimal amount of additional hardware and clock cycles/-control states. Remember when adding new instructions, don't break the operation of the standard ones.
- (c) (1 mark) Discuss the effect of the modification in part (b) on the latency of single-cycled non-pipelined CPU
- (d) (2 marks) Discuss if the suggested modification in part (b) should be handled by increasing/decreasing the number of pipelining stages that were discussed in class. Draw a pipelined version of the new processor similar to Figure 4.49 of the book.
- (e) (2 marks) Discuss if any new types of data hazards are introduced due to the new instruction? If yes, can they be mitigated through forwarding? Use a multi-cycle pipeline diagram like Figure 4.51 of the book to illustrate the new forwarding paths.
- (f) (1 mark) Discuss, based on the above analysis, why the new instruction was not made part of the RISC-V architecture.

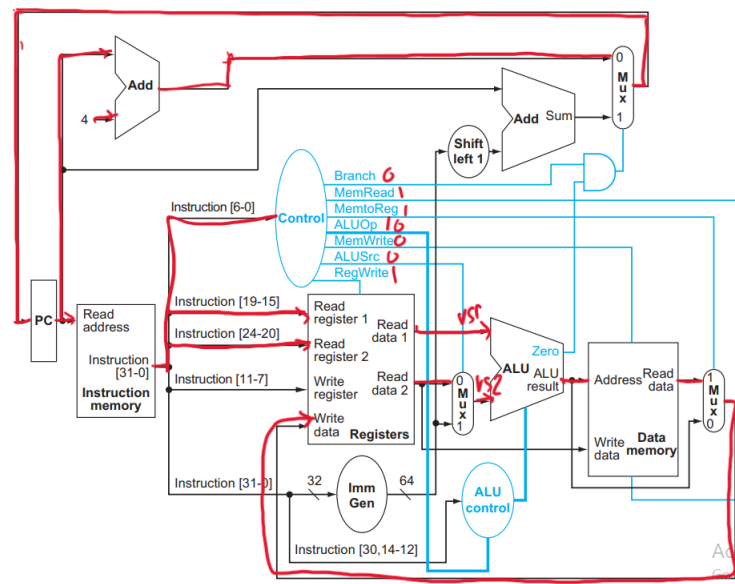
**Solution:****(i) Load Word Register**

```
lwr rd, rs2(rs1)    // Reg[rd] = Mem[Reg[rs1] + Reg[rs2]]
```

- a. We can see that this `lwr` instruction is similar to the R-Type instruction. So we can use the existing R-Type instruction format along with some modification to implement the `lwr` instruction. The R-Type can be made as follows: 7-bit Opcode, 5-bit RD, 3-bit Funct3, 5-bit RS1, 5-bit RS2, 7-bit Funct7
- b. Since the instruction deals with memory read, and register write, the control signals of the traditional R-Type won't work, and would have to be modified:

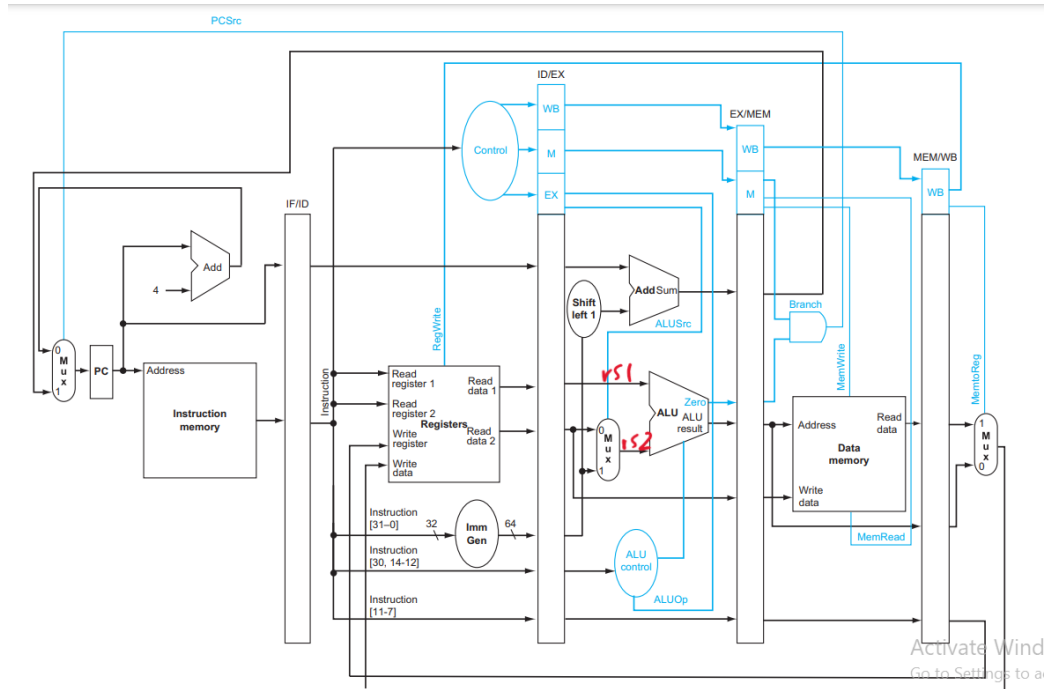
ALUSrc = 0 to add rs2 to rs1 for the offset, ALUOp = 10 (same as add), RegWrite = 1 for write back into RD, MemWrite = 0, MemRead = 1 as memory is read, and MemtoReg = 1 as data from Memory is written on register RD.

The datapath for `lwr` is as follows (red lines show active datapath - not including control signals for which values are written along with the signal):



- c. The latency of single-cycled non-pipelined CPU would remain the same.

d. No changes are needed so the pipelined processor remains same as shown:

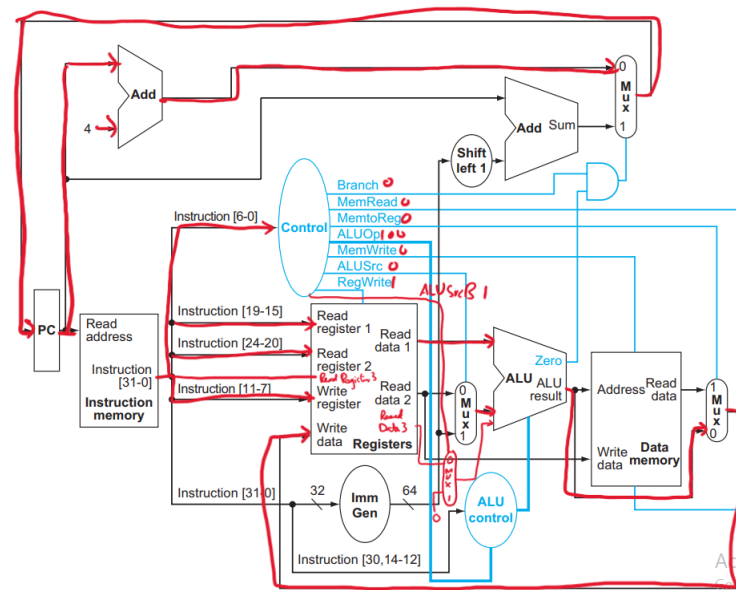


e. no new data hazards would be introduced.

f. RISC-V architecture is the simplest architecture and supports the basic instructions, while lwr is a complex instruction. Hence was not made part of the RISC-V architecture. Also it does not make much sense to add register value as an offset when the ld instruction already exists which basically performs the same task.

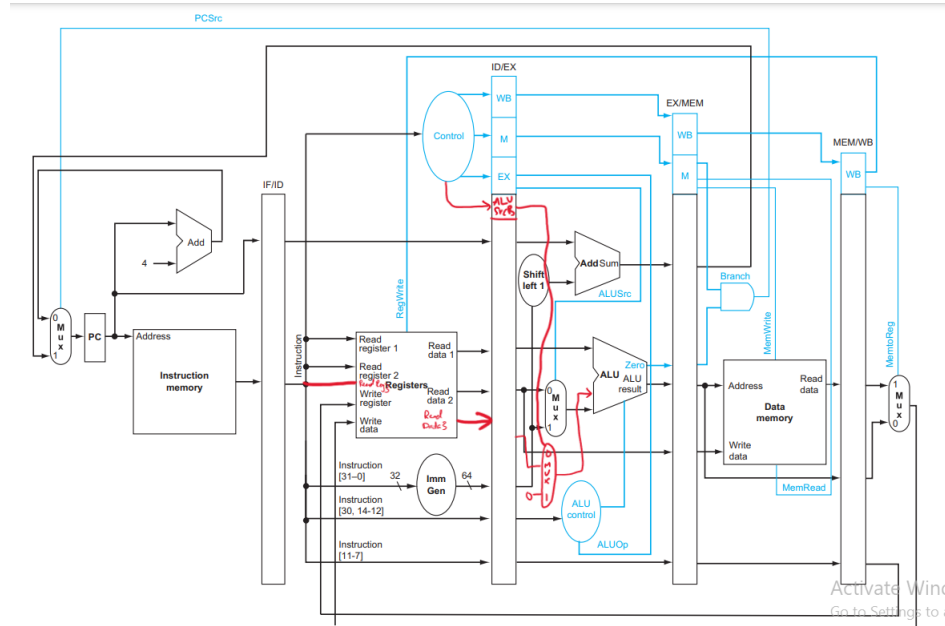
```
add3 rd, rs1, rs2, rs3    // Reg[rd] = Reg[rs1] + Reg[rs2] + Reg[rs3]
```

- The datapath for `add3` is as follows:



- c. It would definitely increase the latency, as now we would require three register reads, three register value outputs, and three ALU operands. So each of the ID stage, and EX stage would then take more time and hence increased latency.
- d. It can be handled with the same number of pipeline stages, however, we would

require an increased number of registers at some stages (ID/EX and EX/MEM) and would have to handle more signals as shown:

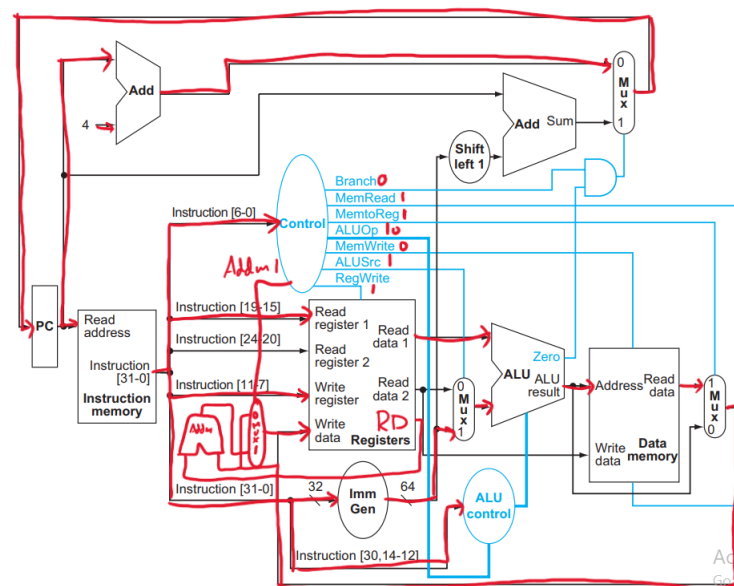


- e. There wouldn't be any new data hazard signals, only the previous ones such as a data hazard which can be mitigated through forwarding.
- f. RISC-V architecture is the simplest architecture. Adding this instruction would increase the complexity, hence it was not added to the RISC-V architecture, as the same instruction can be performed by the use of two add instructions as well.

(iii) Add to Memory


`addm rd, Offset(rs)     //  $\text{Reg}[\text{rd}] = \text{Reg}[\text{rd}] + \text{Mem}[\text{Offset} + \text{Reg}[\text{rs}]]$`

- This instruction deals with both a register operand, and memory access. However, it seems similar to the `ld` instruction format. So we can encode the `addm` instruction using the I-Type instruction format: 7-bit Opcode, 5-bit RD, 3-bit Funct3, 5-bit RS, 12-bit Immediate value.
- Since we are adding the value of the offset to RS, and then further adding it to RD, we can implement an additional MUX to the write back signal that chooses between an adder value of RD generated from the Register File, or the simple write back from memory - a signal “Addm” would be given from the control unit. If the signal is 1, then added value is selected, else the original write back value is selected. The datapath then becomes as shown:



- It would increase the latency as the new instruction requires another signal “Addm” and additional hardware; Mux and Adder for the writeback.
- No new pipeline stages would be needed, but “Addm” signal would be needed till the MEM/WB pipeline after which in the writeback stage we will place the adder and Mux. As shown:





- e. There won't be any new data hazards but the same as before.
- f. RISC-V architecture is the simplest architecture. Adding this instruction would increase the complexity significantly, hence it was not added to the RISC-V architecture. Further, the same operation can be performed using an `ld` and an `add` instruction.

(iv) Branch Equal to Memory

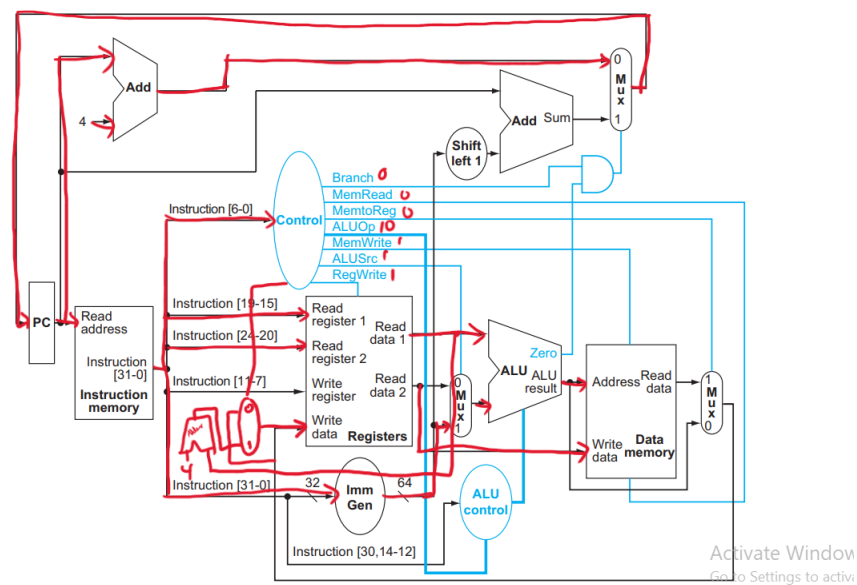
```
beqm rs1, offset(rs2), rs3    // if(Reg[rs1] == Mem[Offset + Reg[rs2]])  
                               //      PC = PC + Reg[rs3]
```

- a. There is no existing RISC-V instruction format that is suitable to encode this instruction as this branch instruction involves a memory access operand in the instruction. So we will have to define a new instruction such as “BM-Type” which denotes a branch instruction with memory access. The instruction format can then be 7-bit Opcode [custom], 5-bit RS1, 5-bit RS2, 5-bit RS3, 3-bit funct3 value [custom], and 7-bit immediate value.
- b. There is no WriteData or writeback, however, we need to access Memory along with offset on RS2, and then we need to perform ALU operation (a - b) to check for 0 value. If zero, then we need to increment PC value with RS3 value. RS3 is an input to Register File, and output is ReadData3 which is sent as an increment to the PC value counter. The datapath becomes:
- c. The latency would undoubtedly increase as now we have to assume branch is taken, then access the memory, then again access the ALU Stage to calculate branch, and in case the branch is true, then send a signal such that PC value is incremented by the value of RS3. So it would require more clock cycles or more time to execute, so latency increases.
- d. We would have to increase the number of pipeline stages to accomodate for the increase in hardware and increased clock cycles to allow the execution of instructions - an additional ALU might be considered.
- e. There won't be any new types of data hazards.
- f. RISC-V architecture is the simplest architecture. Adding this instruction would increase the complexity of the design significantly, as we need more signals, more hardware, and more pipeline stages. Hence it was not added to the RISC-V architecture.

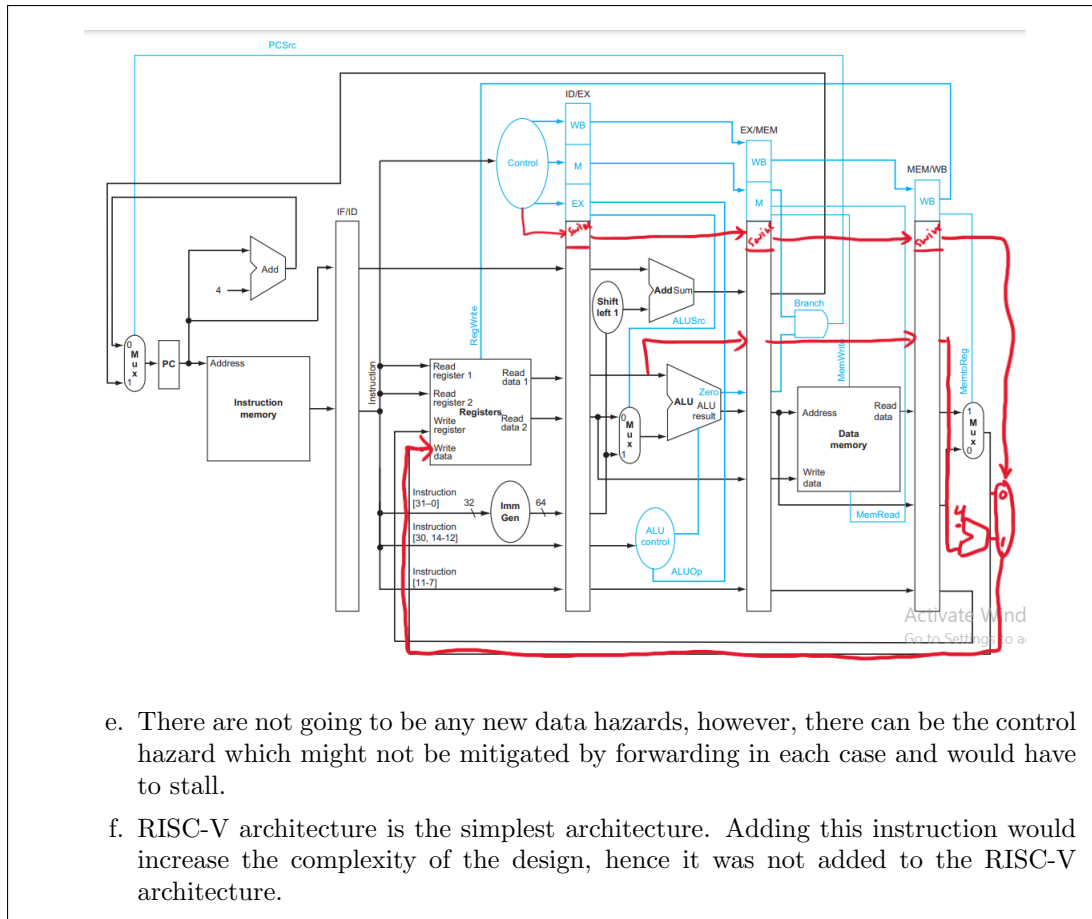
(v) Store Word and Increment

```
swinc rs2, offset(rs1)    // Mem[Reg[rs1] + offset] = Reg[rs2],
                          // Reg[rs1] = Reg[rs1] + 4
```

- There is no existing RISC-V instruction format that can access the memory to store a value and also increment the value in a register simultaneously, however, we can use and modify the existing S-Type instruction format. The instruction format can then be 7-bit Opcode [custom], 5-bit Immediate, 3-bit Funct3 [custom], 5-bit RS1, 5-bit RS2, 7-bit Immediate.
- We would have a MemWrite where the write data is the value of RS2, and at the same time we would increment the value of RS1 by 4 using an adder, and a MUX controlled by a “swinc” signal as below:



- The latency would increase as we would now have to assert signals for accessing memory and incrementing the value of the register through additional hardware which would definitely take more time.
- There is no need for increasing pipeline stages, however, there will be more signals and hardware as shown:



**Question 8: Structural Hazard Analysis**

Consider the fragment of the RISC-V assembly below:

```
ld    x31, 32(x10)
sd    x11, 8(x10)
sub   x12, x14, x13
add   x15, x12, x13
beq   x24, x0, label
add   x31, x31, x14
```

Suppose we modify the pipeline so that it has only one memory (that handles both instructions and data). In this case, there will be a structural hazard every time a program needs to fetch an instruction during the same cycle in which another instruction accesses data.

- (i) Draw a pipeline diagram to show where the code above will stall.

**Solution:** Stall has been highlighted with yellow in the figure below.

Instruction \ Cycle	1	2	3	4	5	6	7	8	9	10	11	12
ld x31, 32(x10)	IF	ID	EX	MEM	WB							
sd x11, 8(x10)		IF	ID	EX	MEM	WB						
sub x12, x14, x13			IF	ID	EX	MEM	WB					
add x15, x12, x13				Stall	Stall	IF	ID	EX	MEM	WB		
beq x24, 0, label							IF	ID	EX	MEM	WB	
add x31, x31, x14								IF	ID	EX	MEM	WB

- (ii) In general, is it possible to reduce the number of stalls/NOPs resulting from this structural hazard by reordering code?

**Solution:** Reordering can sometimes help, however, for this structural hazard, re-ordering the code will not help as every instruction must be fetched; hence we will encounter a stall every time data has to be accessed.

- (iii) Must this structural hazard be handled in hardware? We have seen that data hazards can be eliminated by adding NOPs to the code. Can you do the same with this structural hazard? If so, explain how. If not, explain why not.

**Solution:** NOPs will also have to be fetched from the Instruction Memory, so again in the fetching stage, there will be a stall. So adding NOPs will not resolve the structural hazard.

- (iv) Approximately how many stalls would you expect this structural hazard to generate in a typical program? Use this instruction mix from Question 2.

**Solution:** Since every data access would cause a stall, then stall would occur each time Data Memory is accessed; load and store instructions.  
Then we would expect stall on  $20\% + 25\% = 45\%$  of the instructions using the instruction mix from Question 2.

**Question 9: Forwarding/Hazard-Detection Units Analysis [5 marks]**

Problems in this exercise refer to the following sequence of instructions, and assume that it is executed on a five-stage pipelined datapath.

```
and x16, x10, x9
ld x29, 4(x16)
ld x10, 0(x2)
sub x29, x29, x15
sd x29, 0(x16)
```

- (i) If there is no forwarding or hazard detection, insert NOPs to ensure correct execution.

**Solution:** There will be two stalls after **and** instruction, as the next instruction uses the destination register of **and**. There will be a stall after **ld** instruction, and two stalls after **sub** instruction. Then we have to insert NOPs at these occurrences.

```
and x16, x10, x9
NOP
NOP
ld x29, 4(x16)
ld x10, 0(x2)
NOP
sub x29, x29, x15
NOP
NOP
sd x29, 0(x16)
```

- (ii) Now, change and/or rearrange the code to minimize the number of NOPs needed. You can assume register x17 can be used to hold temporary values in your modified code.

**Solution:** NOPs won't be reduced by reordering the code in this case.

- (iii) If the processor has forwarding, but we forgot to implement the hazard detection unit, what happens when the original code executes?

**Solution:** The code will still execute without any complications as hazard detection is needed in case of a load instruction such that it uses the result of the previous load instruction which is not happening in the given sequence of instructions.

- (iv) If the processor has forwarding, for the first seven cycles during the execution of this code, specify which signals are asserted in each cycle by hazard detection and forwarding units in Figure 4.59 of the book.

**Solution:**

Instruction \ Cycle	1	2	3	4	5	6	7	8	9
and x16, x10, x9	IF	ID	EX	MEM	WB				
ld x29, 4(x16)		IF	ID	EX	MEM	WB			
ld x10, 0(x2)			IF	ID	EX	MEM	WB		
sub x29, x29, x15				IF	ID	EX	MEM	WB	
sd x29, 0(x16)					IF	ID	EX	MEM	WB

There are no stalls for the given sequence of instructions, and there would be no signals asserted by the hazard detection unit.

No execution for the first two cycles, so no forwarding.

(3) First execution, ForwardA = 00, ForwardB = 00

(4) Need for forwarding ALU Result to ld; ForwardA = 10, ForwardB = 00

(5) No forwarding; ForwardA = 00, ForwardB = 00

(6) Need for forwarding result of ld to sub; ForwardA = 00, ForwardB = 01

(7) Need for forwarding result of sub to sd; Forward A = 10, ForwardB = 00.

\*[The above ForwardA and ForwardB values are in binary as from the book]

- (v) If there is no forwarding, what new input and output signals do we need for the hazard detection unit in Figure 4.59? Using this instruction sequence as an example, explain why each signal is needed.

**Solution:** The hazard detection unit would additionally need to check if the result of the previous instruction is being used by the current/next instruction, in which case it would need to forward the result. So the hazard detection unit would need the rd values from the MEM/WB register. If the new instruction in the ID needs the result of the previous instruction, then the current instruction in the ID would be stalled. So the rd values of both the registers would have to be compared. The hazard unit would already have the destination register from the EX/MEM register, so only the value from MEM/WB has to be sent to the hazard detection unit.

So no additional outputs will be needed, and the pipeline can be stalled with the existing signals by just reconnecting them.

Using this instruction sequence as an example, the value of rd from EX/MEM is needed to detect the data hazard between **and** and the **ld** instruction. The value of rd from MEM/WB is needed to detect the data hazard between the first **ld** instruction and **sub** instruction.

- (vi) For the new hazard detection unit form (v), specify which output signals it asserts in each of the first five-cycles during the execution of this code.

**Solution:**

Instruction \ Cycle	1	2	3	4	5
and x16, x10, x9	IF	ID	EX	MEM	WB
ld x29, 4(x16)		IF	ID	-	-
ld x10, 0(x2)			IF	-	-

In the first three cycles, PCWrite = 1, IF/ID Write = 1, and control mux = 0 as there is no data hazard, so no stall or forwarding required and things operate as normal.

(4) PCWrite = 0, IF/ID Write = 0, control mux = 1 as there is a hazard now so we need to stall the instruction.

(5) PCWrite = 0, IF/ID Write = 0, control mux = 1 due to hazard and need for stall.

#### Question 10: Forwarding Trade-Offs [5 marks]

Consider the pipelined RISC-V version with 5 stages discussed in class that does not handle data hazards (i.e., the programmer is responsible for addressing data hazards by inserting NOP instructions where necessary). Suppose that (after optimization) a typical  $n$ -instruction program requires an additional  $0.3 \cdot n$  NOP instructions to correctly handle data hazards.

- (i) Suppose that the cycle time of this pipeline without forwarding is 200ps. Suppose also that adding forwarding hardware will reduce the number of NOPs from  $0.3 \cdot n$  to  $0.06 \cdot n$ , but increase the cycle time to 250ps. What is the speedup of this new pipeline compared to the one without forwarding?

**Solution:** Pipeline without forwarding takes  $1.3 \cdot n \cdot 200ps = 260n$ .

Pipeline with forwarding takes  $1.06 \cdot n \cdot 250ps = 265n$

$$\text{Speedup} = \frac{260}{265} = 0.98$$

- (ii) Different programs will require different amounts of NOPs. How many NOPs (as a percentage of code instructions) can remain in the typical program before that program runs slower on the pipeline with forwarding?

**Solution:** Pipeline with forwarding must be faster than the pipeline without the forwarding. Then  $250 \cdot (1 + s) \cdot n < 200 \cdot 1.3 \cdot n$  where  $s$  is the number of stalls.

$$\text{Then } 1 + s < 1.04 \implies s < 0.04.$$

So  $s$  must be less than 4%.

- (iii) Repeat (ii); however, this time let  $x$  represent the number of NOP instructions relative to  $n$ . (In (ii),  $x$  was equal to 0.3) Your answer will be with respect to  $x$ .



**Solution:** Now we have  $x$  relative to our answer in part (ii) where  $x$  is the number of NOPs relative to  $n$ . Then  $250 * (1 + s) * n < 200 * (1 + x) * n$ .

$$\begin{aligned} \text{Then } 1 + s &< \frac{4(1+x)}{5} \\ \implies s &< \frac{4 + 4x - 5}{5} \implies s < \frac{4x - 1}{5} \end{aligned}$$

- (iv) Can a program with only  $.075 * n$  NOPs (in the no-forwarding case) possibly run faster on the pipeline with forwarding? Explain why or why not.

**Solution:** We have  $0.075 * n$  NOPs. With forwarding, the program takes  $250ps$  while without forwarding, the program takes  $200 * 1.075 = 215ps$ . Hence it will not be able to run faster.

- (v) At minimum, how many NOPs (as a percentage of code instructions) must a program have before it can possibly run faster on the pipeline with forwarding?

**Solution:** We need the program to run faster on the pipeline with forwarding, therefore speedup must be positive.

Then  $\frac{4x - 1}{5} > 0 \implies x > 0.25$ . Then at minimum, the program must have at least 0.25 NOPs to run faster on the pipeline with forwarding.

**Question 11: Forwarding Logic Design Trade-Offs [5 marks]**

This exercise is intended to help you understand the cost/complexity/ performance trade-offs of forwarding in a pipelined processor. Problems in this exercise refer to pipelined datapaths from Figure 4.53 of the book. These problems assume that, of all the instructions executed in a processor, the following fraction of these instructions has a particular type of read-after-write (RAW) data dependence.

EX to 1st Only	MEM to 1st Only	EX to 2nd Only	MEM to 2nd Only	EX to 1st and EX to 2nd
10%	25%	10%	15%	15%

The type of RAW data dependence is identified by the stage that produces the result (EX or MEM) and the next instruction that consumes the result (1st instruction that follows the one that produces the result, 2nd instruction that follows, or both). We assume that the register write is done in the first half of the clock cycle and that register reads are done in the second half of the cycle, so “EX to 3rd” and “MEM to 3rd” dependences are not counted because they cannot result in data hazards. We also assume that branches are resolved in the EX stage (as opposed to the ID stage), and that the CPI of the processor is 1 if there are no data hazards.

Assume the following latencies for individual pipeline stages. For the EX stage, latencies are given separately for a processor without forwarding and for a processor with different kinds of forwarding.

IF	ID	EX (no FW)	EX (full FW)	EX (FW from EX/ MEM only)	EX (FW from MEM/ WB only)	MEM	WB
150ps	120ps	140ps	180ps	150ps	150ps	150ps	120ps

- (i) For each RAW dependency listed above, give a sequence of at least three assembly statements that exhibits that dependency.

**Solution:****(1) EX to 1st Only:**

```
add x5, x6, x7
add x8, x5, x9
add x10, x11, x12
```

**(2) MEM to 1st Only:**

```
ld x10, 0(x12)
add x11, x10, x13
add x5, x6, x7
```

**(3) EX to 2nd Only:**

```
add x5, x6, x7
add x10, x11, x12
add x8, x5, x9
```

**(4) MEM to 2nd Only:**

```
ld    x10, 0(x12)
add   x5, x6, x7
add   x11, x10, x13
```

**(5) EX to 1st and EX to 2nd:**

```
add   x5, x6, x7
add   x8, x5, x9
add   x10, x5, x11
```

- (ii) For each RAW dependancy above, how many NOPs would need to be inserted to allow your code from (i) to run correctly on a pipeline with no forwarding or hazard detection? Show where the NOPs could be inserted.

**Solution:****(1) EX to 1st Only: 2 NOPs**

```
add   x5, x6, x7
NOP
NOP
add   x8, x5, x9
add   x10, x11, x12
```

**(2) MEM to 1st Only: 2 NOPs**

```
ld    x10, 0(x12)
NOP
NOP
add   x11, x10, x13
add   x5, x6, x7
```

**(3) EX to 2nd Only: 1 NOP**

```
add   x5, x6, x7
add   x10, x11, x12
NOP
add   x8, x5, x9
```

**(4) MEM to 2nd Only: 1 NOP**

```
ld    x10, 0(x12)
add   x5, x6, x7
NOP
add   x11, x10, x13
```

**(5) EX to 1st and EX to 2nd: 2 NOPS**

```
add   x5, x6, x7
NOP
NOP
add   x8, x5, x9
add   x10, x5, x11
```

- (iii) Analyzing each instruction independently will over-count the number of NOPs needed to run a program on a pipeline with no forwarding or hazard detection. Write a sequence of three assembly instructions so that, when you consider each instruction in the sequence independently, the sum of the stalls is larger than the number of stalls the sequence actually needs to avoid data hazards.

**Solution:** A possible code can be considered:

```
ld    x10, 0(x12)      # MEM to 2nd: 1 NOP
add   x5,  x6,    x7    # EX to 1st: 2 NOPs
add   x11, x10,    x5
add   x28, x29,    x30
```

Individually if the code is analysed, then we would have 3 NOPs, 1 from the first ld instruction and 2 from the second add instruction. However, when the code will be executed, there will only be 2 NOPs, or 2 stalls occurring only after the second add instruction, as there is overlap between the NOPs resulting from the first ld instruction, and the second add instruction, hence the sum of NOPs from the code is less than the sum of NOPs obtained after analysing each instruction individually.

Like so:

```
ld    x10, 0(x12)      # MEM to 2nd: 1 NOP
add   x5,  x6,    x7    # EX to 1st: 2 NOPs
NOP
NOP
add   x11, x10,    x5
add   x28, x29,    x30
```

- (iv) Assuming no other hazards, what is the CPI for the program described by the table above when run on a pipeline with no forwarding? What percent of cycles are stalls? (For simplicity, assume that all necessary cases are listed above and can be treated independently.)

**Solution:** We can take an average of the codes used in (ii) to get an estimate of stalls per instruction.

Average =  $2 \times 10\% + 2 \times 25\% + 1 \times 10\% + 1 \times 15\% + 2 \times 15\% = 1.25$  stalls per instruction  
on average for a CPI of 2.25. Then on average, there are  $\frac{1.25}{2.25} = 0.556$  cycles or 55.6% stalls.

- (v) What is the CPI if we use full forwarding (forward all results that can be forwarded)? What percent of cycles are stalls?

**Solution:** The only dependency that can not be handled by forwarding is the MEM to 1st dependency - forwarding from MEM stage to the next instruction. Then 25% of instructions will generate 1 stall. Then there are  $\frac{0.25}{1.25} = 0.2$  cycles or 20% stalls.

- (vi) Let us assume that we cannot afford to have three-input multiplexors that are needed for full forwarding. We have to decide if it is better to forward only from the EX/MEM pipeline register (next-cycle forwarding) or only from the MEM/WB pipeline register (two-cycle forwarding). What is the CPI for each option?

**Solution:**

(1) **From EX/MEM:**

Stalls from forwarding from EX/MEM are 0 in EX to 1st, 2 in MEM to 1st, 1 in EX to 2nd, 1 in MEM to 2nd, and 1 in EX to 1st and 2nd. Then we get an average of  $0 \times 10\% + 2 \times 25\% + 1 \times 10\% + 1 \times 15\% + 1 \times 15\% = 0.9$  stalls per instruction. **CPI = 1.9.**

(2) **From MEM/WB:**

Stalls are 1 in EX to 1st, 1 in MEM to 1st, 0 in EX to 2nd, 0 in MEM to 2nd, 1 in EX to 1st and 2nd. Then on average  $1 \times 10\% + 1 \times 25\% + 0 \times 10\% + 0 \times 15\% + 1 \times 15\% = 0.5$  stalls per instruction. **CPI = 1.5.**

- (vii) For the given hazard probabilities and pipeline stage latencies, what is the speedup achieved by each type of forwarding (EX/MEM, MEM/WB, for full) as compared to a pipeline that has no forwarding?

**Solution:** The following table shows the necessary information required to calculate the relative speedup of each type of forwarding as compared to a pipeline without forwarding.

	No Forwarding	EX/MEM	MEM/WB	Full Forwarding
CPI	2.25	1.9	1.5	1.2
Clock Rate	150	150	150	180
Time	$2.25 \times 150 = 337.5$	$1.9 \times 150 = 285$	$1.5 \times 150 = 225$	$1.2 \times 180 = 216$
Speedup	-	1.18	1.5	1.56

- (viii) What would be the additional speedup (relative to the fastest processor from (vii)) be if we added “time-travel” forwarding that eliminates all data hazards? Assume that the yet-to-be-invented time-travel circuitry adds 100ps to the latency of the full-forwarding EX stage.

**Solution:**

MEM/WB: CPI = 1.5, Clock Rate = 150, Time = 225

Time Travel: CPI = 1.00, Clock Rate = 280, Time = 280.0

$$\text{Speedup by time travel} = \frac{225}{280} = 0.80$$

The additional speedup relative to the fastest processor from vii is of 0.80 which means that time-travel would slow the processor further if it adds 100ps to the latency of the full forwarding.

- (ix) The table of hazard types has separate entries for “EX to 1st” and “EX to 1st and EX to 2nd”. Why is there no entry for “MEM to 1st and MEM to 2nd”?

**Solution:** From part (vi), we can see that in the EX/MEM stage, there is no stall from EX to 1st, but there is one stall from EX to 1st and 2nd. However, MEM to 1st and MEM to 1st and 2nd both will have the same number of stalls - all MEM to 1st will cause a stall due to which the second instruction's ID overlaps with the original instruction's WB. Hence there is no entry for MEM to 1st and MEM to 2nd.